

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [19]: import warnings
warnings.filterwarnings('ignore')
```

```
In [8]: data=pd.DataFrame(pd.read_csv("advertising.csv"))
data
```

```
Out[8]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [9]: data.head()
```

```
Out[9]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [11]: data.head(11)
```

```
Out[11]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
5	8.7	48.9	75.0	7.2
6	57.5	32.8	23.5	11.8
7	120.2	19.6	11.6	13.2
8	8.6	2.1	1.0	4.8
9	199.8	2.6	21.2	15.6
10	66.1	5.8	24.2	12.6

```
In [12]: data.shape
```

```
Out[12]: (200, 4)
```

```
In [13]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype  
---  -
0    TV          200 non-null   float64
1    Radio       200 non-null   float64
2    Newspaper   200 non-null   float64
3    Sales       200 non-null   float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [14]: data.describe()
```

```
Out[14]:
```

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

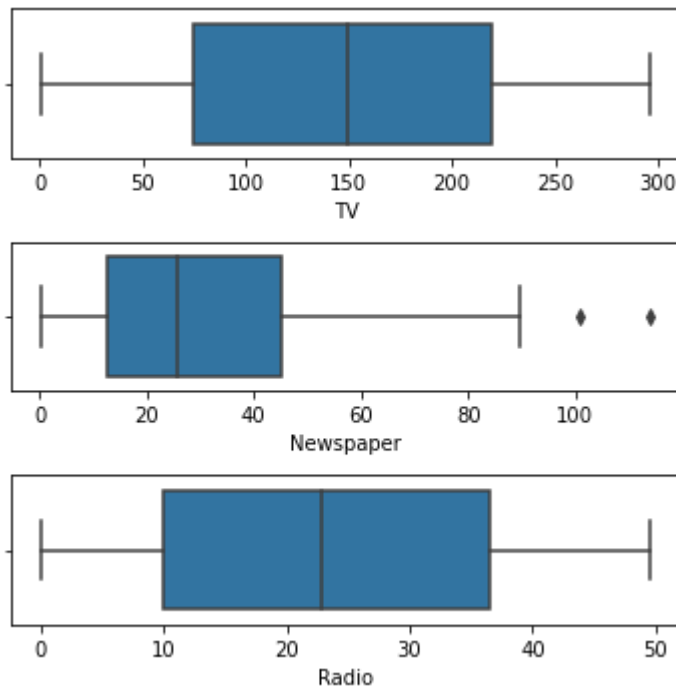
```
In [16]: data.isnull().sum()
```

```
Out[16]: TV          0  
Radio         0  
Newspaper     0  
Sales         0  
dtype: int64
```


```
In [17]: data.isnull().sum()*100/data.shape[0]
```

```
Out[17]: TV          0.0  
Radio         0.0  
Newspaper     0.0  
Sales         0.0  
dtype: float64
```


```
In [20]: fig, axs = plt.subplots(3, figsize = (5,5))  
plt1 = sns.boxplot(data['TV'], ax = axs[0])  
plt2 = sns.boxplot(data['Newspaper'], ax = axs[1])  
plt3 = sns.boxplot(data['Radio'], ax = axs[2])  
plt.tight_layout()
```




```
In [22]: sns.boxplot(data['Sales'])  
plt.show()
```



```
In [31]: #Let's see how Sales are related with other variables using scatter plot.
sns.pairplot(data, x_vars=['TV', 'Newspaper', 'Radio'], y_vars='Sales', height=4, a
plt.show())
```



```
In [33]: sns.heatmap(data.corr(), cmap="YlGnBu", annot = True)
plt.show()
#As is visible from the pairplot and the heatmap, the variable TV seems to be most
#So Let's go ahead and perform simple linear regression using TV as our feature var
```



```
In [34]: X = data['TV']
y = data['Sales']
```

```
In [35]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7, test_si
```

```
In [37]: X_train.head()
```

```
Out[37]: 74      213.4
3        151.5
185      205.0
26       142.9
90       134.3
Name: TV, dtype: float64
```

```
In [38]: y_train.head()
```

```
Out[38]: 74      17.0
3        16.5
185      22.6
26       15.0
90       14.0
Name: Sales, dtype: float64
```

```
In [39]: import statsmodels.api as sm
```

```
In [ ]: #By default, the statsmodels library fits a line on the dataset which passes throug
#But in order to have an intercept, you need to manually use the add_constant attri
#And once you've added the constant to your X_train dataset, you can go ahead and f
#line using the OLS (Ordinary Least Squares) attribute of statsmodels as shown belo
```

```
In [42]: # Add a constant to get an intercept
X_train_sm = sm.add_constant(X_train)

# Fit the resgression line using 'OLS'
lr = sm.OLS(y_train, X_train_sm).fit()
```

```
In [43]: # Print the parameters, i.e. the intercept and the slope of the regression line fit
lr.params
```

```
Out[43]: const      6.948683
TV          0.054546
dtype: float64
```

```
In [ ]:
```

```
In [44]: # Performing a summary operation lists out all the different parameters of the regr
print(lr.summary())
```

OLS Regression Results


=====						
Dep. Variable:	Sales	R-squared:	0.816			
Model:	OLS	Adj. R-squared:	0.814			
Method:	Least Squares	F-statistic:	611.2			
Date:	Mon, 06 Mar 2023	Prob (F-statistic):	1.52e-52			
Time:	10:50:37	Log-Likelihood:	-321.12			
No. Observations:	140	AIC:	646.2			
Df Residuals:	138	BIC:	652.1			
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	6.9487	0.385	18.068	0.000	6.188	7.709
TV	0.0545	0.002	24.722	0.000	0.050	0.059
=====						
Omnibus:	0.027	Durbin-Watson:		2.196		
Prob(Omnibus):	0.987	Jarque-Bera (JB):		0.150		
Skew:	-0.006	Prob(JB):		0.928		
Kurtosis:	2.840	Cond. No.		328.		
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [49]: plt.scatter(X_train, y_train)
plt.plot(X_train, 6.948 + 0.054*X_train, 'r')
plt.show()
```



```
In [50]: # Add a constant to X_test
X_test_sm = sm.add_constant(X_test)

# Predict the y values corresponding to X_test_sm
y_pred = lr.predict(X_test_sm)
```

```
In [51]: y_pred.head()
```

```
Out[51]: 126      7.374140
          104     19.941482
          99     14.323269
          92     18.823294
          111    20.132392
          dtype: float64
```

```
In [52]: from sklearn.metrics import mean_squared_error
          from sklearn.metrics import r2_score
```


```
In [53]: #Returns the mean squared error; we'll take a square root
          np.sqrt(mean_squared_error(y_test, y_pred))
```

```
Out[53]: 2.019296008966233
```

```
In [54]: r_squared = r2_score(y_test, y_pred)
          r_squared
```

```
Out[54]: 0.7921031601245658
```

```
In [55]: plt.scatter(X_test, y_test)
          plt.plot(X_test, 6.948 + 0.054 * X_test, 'r')
          plt.show()
```



```
In [ ]:
```